


AgenticTrip — Phase 1 Architecture & Execution Plan

Project Vision

AgenticTrip is an AI-powered travel planning and execution platform that uses autonomous AI agents, MCP integrations, real-time APIs, memory, and conversational workflows to:

- Understand user travel goals naturally
- Plan complete trips end-to-end
- Book and manage travel services
- Provide real-time recommendations
- Adapt dynamically during travel
- Support multilingual interaction
- Operate as a multi-agent ecosystem

High-Level Architecture



Core Technology Stack

Frontend

- React + Vite
- TailwindCSS
- ShadCN UI
- Framer Motion
- Zustand or Redux Toolkit

Backend

- FastAPI
- Python 3.11+
- Uvicorn
- Pydantic

AI / Agent Framework

Recommended:

Option 1 (Recommended)

- LangGraph
- LangChain
- Azure OpenAI

Option 2

- CrewAI

Option 3

- Semantic Kernel

Recommended AI Models

Primary LLM

- Azure OpenAI GPT-4.1

Embedding Model

- text-embedding-3-large

Speech Models (Future)

- Whisper
 - Azure Speech
-

Phase-wise Development Plan

PHASE 1 — Foundation Setup

Goal

Create production-ready base architecture.

Deliverables

1. Frontend Setup

```
npm create vite@latest agentictrip
```

Install:

```
npm install tailwindcss react-router-dom axios  
npm install framer-motion lucide-react
```

2. Backend Setup

```
mkdir backend  
cd backend  
python -m venv venv
```

Install:

```
pip install fastapi uvicorn python-dotenv  
pip install openai langchain langgraph  
pip install chromadb faiss-cpu  
pip install tiktoken pydantic
```

3. Environment Configuration

Create `.env`

```
AZURE_OPENAI_API_KEY=  
AZURE_OPENAI_ENDPOINT=  
AZURE_OPENAI_DEPLOYMENT=  
AZURE_OPENAI_API_VERSION=
```

4. Basic FastAPI Structure

```
backend/  
|  
├─ app/  
|   ├── agents/  
|   ├── tools/  
|   ├── memory/  
|   ├── rag/  
|   ├── api/  
|   ├── services/  
|   └─ main.py  
├─ requirements.txt  
└─ .env
```

PHASE 2 — Multi-Agent System

Initial Agents

1. Planner Agent

Responsibilities:

- Understand travel intent
- Build itinerary
- Coordinate other agents

2. Budget Agent

Responsibilities:

- Estimate trip cost
- Suggest optimizations
- Currency conversion

3. Destination Agent

Responsibilities:

- Suggest attractions
- Recommend local experiences
- Generate travel insights

4. Booking Agent

Responsibilities:

- Flights
- Hotels
- Transportation

5. Weather Agent

Responsibilities:

- Forecast analysis
- Weather-based recommendations

6. Support Agent

Responsibilities:

- FAQs
- Emergency support
- Trip modifications

PHASE 3 — MCP Integration Layer

MCP Concept

MCP (Model Context Protocol) allows AI agents to use external tools/services in a standardized way.

MCP Integrations Needed

Capability	MCP / API
Maps	Google Maps API
Hotels	Booking.com API
Flights	Amadeus API
Weather	OpenWeatherMap API
Currency	ExchangeRate API
Search	Tavily / SerpAPI
Calendar	Google Calendar API
Email	Gmail API
Payments	Stripe

PHASE 4 — RAG + Memory

Goal

Allow agents to remember users and retrieve contextual travel data.

Long-Term Memory

Store:

- User preferences
 - Past trips
 - Budget patterns
 - Favorite destinations
 - Language preference
-

RAG Sources

Possible datasets:

- City guides
- Tourism PDFs

- Hotel reviews
 - Airline policies
 - Visa rules
 - User uploaded documents
-

Vector DB Options

Recommended:

- ChromaDB
- FAISS

Future:

- Pinecone
 - Weaviate
-

PHASE 5 — Real-Time Agent Workflows

Example Workflow

User:

Plan a 5-day trip to Bali under ₹80,000

Agent Flow

```
Planner Agent
  ↓
Budget Agent
  ↓
Flight Search Tool
  ↓
Hotel Search Tool
  ↓
Weather Agent
  ↓
Final Travel Plan
```

PHASE 6 — UI/UX Features

Features

- Chat-based interface
 - Dynamic itinerary cards
 - Map visualization
 - Expense tracker
 - Real-time alerts
 - Voice interaction
 - Multi-language support
 - AI travel assistant sidebar
-

Suggested Database

PostgreSQL

Tables:

- users
 - trips
 - itineraries
 - expenses
 - memories
 - bookings
 - conversations
-

Authentication

Recommended:

- Clerk OR
 - Firebase Auth OR
 - Auth0
-

Deployment Strategy

Frontend

- Azure Static Web Apps OR
- Vercel

Backend

- Azure App Service
- Azure Container Apps
- Dockerized deployment

Database

- Azure PostgreSQL

Vector DB

- Chroma local initially
- Pinecone later

Recommended Repository Structure

```
AgenticTrip/  
|  
├─ frontend/  
├─ backend/  
├─ docs/  
├─ architecture/  
├─ prompts/  
├─ datasets/  
└─ deployment/
```

Immediate Action Items

STEP 1

Create GitHub repository:

AgenticTrip

STEP 2

Setup frontend + backend base project.

STEP 3

Create first Planner Agent.

STEP 4

Integrate Azure OpenAI.

STEP 5

Build simple conversational trip planner.

First Milestone Goal

By the end of Milestone 1, AgenticTrip should:

✓ Accept user trip requests ✓ Understand travel goals ✓ Generate travel itineraries ✓ Use at least 2 tools/APIs ✓ Maintain conversation memory ✓ Support multi-agent orchestration

Suggested Development Sequence

1. Base Architecture

2. Planner Agent
 3. Tool Calling
 4. Memory
 5. RAG
 6. Multi-Agent Coordination
 7. Real-Time APIs
 8. Deployment
 9. Voice Support
 10. Autonomous Workflows
-

Long-Term Vision

Future enhancements:

- Voice-enabled AI travel companion
 - Autonomous booking approvals
 - Real-time flight disruption handling
 - AI-generated travel videos
 - Group trip collaboration
 - Expense splitting
 - AR/VR travel previews
 - AI concierge mode
 - WhatsApp/Telegram integration
-

FastAPI Foundation Learning Notes

What We Built

We created the first working backend service for AgenticTrip using FastAPI and Uvicorn.

Current architecture:

```
graph TD; Browser --> FastAPI_Backend[FastAPI Backend]; FastAPI_Backend --> Python_Logic[Python Logic];
```

Future architecture:

```
React Frontend
  ↓
FastAPI Backend
  ↓
AI Agents
  ↓
Azure OpenAI + APIs + Databases
```

Understanding main.py

```
from fastapi import FastAPI

app = FastAPI(
    title="AgenticTrip API",
    version="1.0.0"
)

@app.get("/")
def home():
    return {
        "message": "Welcome to AgenticTrip"
    }
```

1. Importing FastAPI

```
from fastapi import FastAPI
```

This imports the FastAPI framework into the project.

Purpose:

- Use FastAPI features
 - Create APIs
 - Handle requests and responses
 - Build backend services
-

2. Creating FastAPI Application

```
app = FastAPI(  
    title="AgenticTrip API",  
    version="1.0.0"  
)
```

This creates the backend application instance.

Metadata:

- title → API name
- version → API version

These appear automatically inside Swagger documentation.

3. Creating API Endpoint

```
@app.get("/")
```

This defines an API route.

Meaning:

When someone sends a GET request to `/`, execute the function below.

Examples of future endpoints:

Endpoint	Purpose
<code>/</code>	Home
<code>/chat</code>	AI chat
<code>/trips</code>	Trip management
<code>/hotels</code>	Hotel search
<code>/weather</code>	Weather info

4. Route Function

```
def home():
```

This function runs when the endpoint is accessed.

5. Returning JSON Response

```
return {  
    "message": "Welcome to AgenticTrip"  
}
```

FastAPI automatically converts Python dictionaries into JSON responses.

Browser output:

```
{  
  "message": "Welcome to AgenticTrip"  
}
```

Understanding Uvicorn Command

Command used:

```
uvicorn app.main:app --reload
```

Breakdown

uvicorn

Runs the backend server.

app.main

Meaning:

- Open `app/` folder
 - Open `main.py`
-

:app

Use the FastAPI object named `app`.

--reload

Automatically reloads the server whenever code changes.

Useful during development.

Swagger Documentation

URL:

`http://127.0.0.1:8000/docs`

FastAPI automatically generates:

- API documentation
- Endpoint testing interface
- OpenAPI schema
- Request/response visualization

Benefits:

- Easier frontend integration
 - Better developer experience
 - Enterprise-grade API management
 - Faster testing
-

What FastAPI Automatically Provided

Without writing additional logic, FastAPI already gave:

✓ Running backend server ✓ API routing ✓ JSON handling ✓ Swagger documentation ✓ OpenAPI schema generation ✓ Validation support ✓ Scalable backend architecture

Current AgenticTrip Backend Flow

Current:

User → FastAPI → JSON Response

Future:

```
User
↓
React Frontend
↓
FastAPI Backend
↓
Planner Agent
↓
Azure OpenAI
↓
Travel APIs
↓
AI Response
```

Importance Of This Step

This was the first foundational backend service for AgenticTrip.

This structure will later support:

- AI agents
- Multi-agent orchestration
- Tool calling
- Azure OpenAI integration

- Memory systems
- RAG pipelines
- Real-time APIs
- Production deployment on Azure App Service

Frontend ↔ Backend Integration Notes

Objective

Connect React frontend with FastAPI backend using Axios.

Goal:

Frontend → Backend → Response → UI Update

Architecture After Integration

```
graph TD;
  A[React Frontend] --> B[Axios Request];
  B --> C[FastAPI Backend];
  C --> D[JSON Response];
```

Future architecture:

```
graph TD;
  A[Frontend] --> B[FastAPI];
  B --> C[Planner Agent];
  C --> D[LangGraph];
  D --> E[Azure OpenAI];
  E --> F[Travel APIs];
```

↓
Memory + RAG

Axios Installation

Installed Axios inside frontend.

Command:

```
npm install axios
```

Purpose:

- Send API requests
- Communicate with backend
- Receive JSON responses

API Service Layer

Created:

```
src/services/api.js
```

Code:

```
import axios from "axios"

const api = axios.create({
  baseURL: "http://127.0.0.1:8000"
})

export default api
```

Purpose:

- Centralized API configuration
- Reusable backend communication layer
- Cleaner frontend architecture

Backend Chat Endpoint

Updated:

```
backend/app/main.py
```

Added:

```
from fastapi import FastAPI
from pydantic import BaseModel
from fastapi.middleware.cors import CORSMiddleware
```

CORS Middleware

Added:

```
app.add_middleware(
    CORSMiddleware,
    allow_origins=["*"],
    allow_credentials=True,
    allow_methods=["*"],
    allow_headers=["*"],
)
```

Purpose:

- Allow frontend and backend communication
- Resolve browser security restrictions
- Enable cross-origin requests

Understanding CORS

Frontend:

```
http://localhost:5173
```

Backend:

```
http://127.0.0.1:8000
```

Since ports differ, browser treats them as different origins.

Without CORS:

```
OPTIONS /chat 405 Method Not Allowed
```

appears.

CORS middleware resolved this issue.

Pydantic Request Model

```
class ChatRequest(BaseModel):  
    message: str
```

Purpose:

- Validate frontend request data
 - Ensure proper structure
 - Prevent invalid requests
-

Chat Endpoint

```
@app.post("/chat")  
def chat(request: ChatRequest):  
  
    user_message = request.message  
  
    return {  
        "response": f"AgenticTrip AI received: {user_message}"  
    }
```

Purpose:

- Receive frontend messages
 - Process request data
 - Return backend response
-

Frontend Home.jsx Integration

Implemented:

- React state management
- Axios API calls
- Dynamic UI updates
- Response rendering

Key concepts used:

React State

```
const [message, setMessage] = useState("")
const [response, setResponse] = useState("")
```

Purpose:

- Store user input
 - Store backend response
 - Automatically re-render UI
-

API Request Function

```
const handleSubmit = async () => {

  const res = await api.post("/chat", {
    message: message
  })

  setResponse(res.data.response)
}
```

Purpose:

- Send request to backend
- Receive response
- Update UI dynamically

Full Request Flow

When user clicks:

Plan Trip

This flow occurs:

```
graph TD; A[React UI] --> B[Axios POST Request]; B --> C[FastAPI /chat endpoint]; C --> D[Pydantic validation]; D --> E[Backend processing]; E --> F[JSON response]; F --> G[React state update]; G --> H[UI re-render];
```

Result Achieved

Successfully implemented:

✅ React frontend ✅ FastAPI backend ✅ Axios integration ✅ API communication ✅ POST request handling ✅ Pydantic validation ✅ CORS handling ✅ Dynamic UI rendering ✅ Full-stack workflow foundation

Current AgenticTrip System

```
Frontend (React)
  ↓
Axios API Layer
  ↓
FastAPI Backend
  ↓
Business Logic
```

Importance Of This Milestone

This is the foundational architecture for:

- AI chat systems
- AI copilots
- Agentic workflows
- SaaS platforms
- Enterprise AI applications

This same request pipeline will later support:

- Azure OpenAI
- LangGraph agents
- Travel APIs
- Memory systems
- RAG pipelines
- Autonomous workflows

UI Components Built So Far

Frontend currently contains:

```
components /
└─ Navbar.jsx
```

Pages:

```
pages/  
└─ Home.jsx
```

Layouts:

```
layouts/  
└─ MainLayout.jsx
```

Current UI Features

Implemented:

✓ Premium dark theme ✓ Navbar ✓ Hero section ✓ AI prompt box ✓ Responsive centered layout ✓
Dynamic response card

Next Recommended Tasks

Upcoming frontend improvements:

- gradients
- glow effects
- loading animations
- typing indicator
- glassmorphism
- responsive optimization
- chat history
- better cards
- Enter-key submission
- error handling

Upcoming backend improvements:

- Azure OpenAI integration
 - Planner Agent
 - LangGraph orchestration
 - memory system
 - RAG architecture
-

Session Stopping Point

AgenticTrip is now officially a working full-stack application.

Current status:

✅ Frontend working ✅ Backend working ✅ API communication working ✅ Full-stack architecture established

Next session will continue from UI enhancement and AI integration phase.

Understanding API Requests And FastAPI Flow

Big Picture Architecture

When a user types:

Plan Bali trip

The following workflow occurs:

Frontend
↓
API Request
↓
FastAPI Backend
↓
Azure OpenAI
↓
AI Response
↓
Frontend UI

What Is An API Request?

An API request means:

Frontend asking backend to perform some operation.

The React frontend cannot directly:

- securely call Azure OpenAI
- expose API keys
- manage AI orchestration
- handle backend business logic
- manage memory systems

So frontend communicates with backend through API requests.

Real World Analogy

Frontend = Customer
FastAPI = Waiter
Azure OpenAI = Chef

Workflow:

```
Customer gives order
      ↓
Waiter takes order
      ↓
Chef prepares food
      ↓
Waiter returns response
```

React Frontend Request

Frontend code:

```
const res = await api.post("/chat", {
  message: currentMessage,
  history: messages
})
```

Purpose:

Send request to backend /chat endpoint.

JSON Payload Sent To Backend

Frontend sends:

```
{  
  "message": "Plan Bali trip",  
  "history": []  
}
```

This data travels over HTTP.

What Is HTTP?

HTTP is the communication protocol used between:

- frontend
- backend
- APIs
- servers

Applications communicate using HTTP requests and responses.

Role Of FastAPI

FastAPI acts as the backend controller/orchestrator.

Responsibilities:

- receive requests
 - validate data
 - execute backend logic
 - call Azure OpenAI
 - return responses
-

FastAPI Endpoint

```
@app.post("/chat")
```

Meaning:

```
If frontend sends POST request to /chat,  
execute this function.
```

Receiving Frontend Data

FastAPI receives request here:

```
def chat(request: ChatRequest):
```

Pydantic Validation

```
class ChatRequest(BaseModel):
```

Pydantic validates:

- whether message exists
- whether message is string
- whether history structure is correct

Benefits:

- request safety
- automatic validation
- structured backend architecture

Backend AI Logic

FastAPI then executes:

```
ai_response = generate_ai_response(  
    request.message,  
    request.history  
)
```

Purpose:

- send user message to Azure OpenAI
- include conversation history
- generate contextual AI response

Azure OpenAI Processing

GPT model receives:

- system prompt
- conversation history
- latest user message

and generates:

- itinerary
- travel suggestions
- contextual response

FastAPI Returning Response

```
return {  
    "response": ai_response  
}
```

FastAPI automatically converts Python dictionaries into JSON.

Returned JSON:

```
{  
  "response": "Day 1: Arrival in Bali..."  
}
```

React Receives Response

Frontend receives backend response and updates state:

```
setMessages((prev) => [...prev, aiMessage])
```

This updates UI dynamically.

Complete End-To-End Flow

```
User types message
      ↓
React state updates
      ↓
Axios sends API request
      ↓
FastAPI receives request
      ↓
Pydantic validates data
      ↓
Backend logic executes
      ↓
Azure OpenAI called
      ↓
GPT generates response
      ↓
FastAPI returns JSON
      ↓
React receives response
      ↓
UI updates automatically
```

Why FastAPI Is Important

Without backend:

❌ API keys exposed ❌ no business logic ❌ no memory orchestration ❌ no agent workflows ❌ insecure architecture

FastAPI Enables

✔ Security ✔ AI orchestration ✔ Memory systems ✔ AI agents ✔ RAG pipelines ✔ Databases ✔ Tool calling ✔ Authentication ✔ Scalable architecture

Why Frontend Should NOT Call GPT Directly

Frontend code is visible inside browser.

If API keys are placed in frontend:

- anyone can steal them
- unauthorized usage can occur
- billing costs can increase
- application becomes insecure

Backend protects all sensitive logic.

Real AI SaaS Architecture

```
graph TD;
  A[Frontend (UI Layer)] --> B[FastAPI (Backend Brain)];
  B --> C[AI Services + Agents + Databases];
```

This architecture is used in:

- ChatGPT-style systems
 - AI copilots
 - enterprise AI applications
 - autonomous agent platforms
 - modern AI SaaS products
-

Planner Agent Architecture Upgrade

Objective

Transform AgenticTrip from a generic AI chatbot into a specialized AI Travel Planning Agent.

Previous Architecture

```
User
↓
GPT
↓
Response
```

Limitations:

- generic responses
- inconsistent structure
- no planning workflow
- limited domain specialization

New Planner Agent Architecture

```
User
↓
Planner Agent
↓
Trip Analysis
↓
Budget Analysis
↓
Destination Logic
↓
Itinerary Generation
↓
Final Structured Response
```

This architecture introduces:

- structured reasoning
- travel-focused behavior
- planning workflows
- domain specialization

What Is An AI Agent?

An AI Agent is:

LLM + Reasoning + Workflow + Tools

Instead of only generating text, agents:

- analyze goals
- follow workflows
- perform reasoning
- use tools/APIs
- generate structured outcomes

Phase 1 — Prompt-Orchestrated Planner Agent

Initial implementation uses:

- structured system prompt
- enforced planning behavior
- travel specialization
- response formatting rules

Before introducing:

- LangGraph
- tool calling
- multi-agent orchestration
- autonomous workflows

Planner Agent System Prompt

Updated system prompt:

```
"""
You are AgenticTrip AI,
an expert AI Travel Planning Agent.

Your responsibilities:
```

```
1. Understand user's:
  - destination
  - budget
  - duration
  - travel style
  - activities
  - preferences

2. Generate:
  - day-wise itinerary
  - estimated budget
  - transport suggestions
  - accommodation guidance
  - food recommendations
  - activity planning

3. Always provide:
  - structured markdown output
  - headings
  - bullet points
  - cost estimates
  - optimized travel flow

4. Behave like:
  - premium AI travel concierge
  - intelligent trip planner
  - travel expert

5. If user request is incomplete:
  politely ask follow-up questions.

6. Keep responses:
  - clear
  - premium
  - organized
  - realistic

"""
```

Impact Of Planner Prompt

Before:

- generic chatbot behavior
- inconsistent responses

- weak planning structure

After:

- structured itineraries
- premium travel recommendations
- budget reasoning
- transport planning
- accommodation guidance
- realistic trip organization

Key AI Engineering Concept

The system prompt acts as:

AI personality + rules + reasoning style + behavior control

It strongly influences:

- response quality
- specialization
- formatting
- planning consistency
- reasoning behavior

Current AgenticTrip Architecture

Frontend

- └─ React
- └─ TailwindCSS
- └─ Streaming UI
- └─ Markdown Rendering
- └─ Conversational Chat

Backend

- └─ FastAPI
- └─ Pydantic
- └─ StreamingResponse
- └─ Conversation Memory
- └─ Planner Agent Logic

AI Layer

- └─ Azure OpenAI
- └─ GPT-4o-mini
- └─ Prompt-Orchestrated Planner Agent

Major Capabilities Achieved

AgenticTrip now supports:

✓ Conversational memory ✓ Streaming AI responses ✓ Structured itinerary generation ✓ Travel-specialized reasoning ✓ Markdown-based rendering ✓ Context-aware conversations ✓ Premium travel planning UX

Important Learning

The project has evolved from:

Calling GPT directly

into:

Designing AI behavior intentionally

This is a major distinction between:

- AI user vs
 - AI engineer
-

Next Evolution — Tool-Based Agent Architecture

Current limitation:

GPT still:

- estimates pricing
- guesses recommendations

- lacks real-time external data
-

Goal Of Tool-Based Architecture

Enable AI to use:

- weather APIs
 - flight APIs
 - hotel APIs
 - maps
 - budget calculators
-

Future Agent Architecture

```
graph TD;
  User --> PlannerAgent[Planner Agent];
  PlannerAgent --> ToolDecision[Tool Decision];
  ToolDecision --> ExternalAPIs[External APIs];
  ExternalAPIs --> Reasoning;
  Reasoning --> FinalResponse[Final Intelligent Response];
```

This becomes:

- autonomous agent behavior
 - MCP-style orchestration
 - LangGraph workflows
 - enterprise-grade AI architecture
-

Recommended Next Upgrade

Before integrating external tools, next architectural improvement should be:

Structured JSON Planner Output

Instead of only:

Plain markdown output

Generate:

```
{
  "destination": "...",
  "budget": "...",
  "days": [
    ...
  ]
}
```

Benefits:

- reusable frontend rendering
- itinerary cards
- database storage
- PDF generation
- dashboards
- exports
- trip sharing
- scalable enterprise workflows

Recommended Next Task

Build the foundational FastAPI backend with:

- Azure OpenAI connection
- Planner Agent
- Tool architecture
- Simple travel chat endpoint
- Memory layer

This should be the next implementation step.